

DEVANA PERUPURAYIL

San Diego, CA | (619) 953-7200 | devanaperupurayil@gmail.com | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

PROFESSIONAL SUMMARY

Data Scientist with an M.S. in Data Science from UC San Diego and 2+ years of software engineering experience applying Python, SQL, statistical modeling, and machine learning to forecasting and large-scale data. Built retailer- and item-level forecasts using XGBoost, Random Forest, ARIMA, and Holt-Winters, improving accuracy from 47% to 71% and reducing MAPE by 30%. Experienced in data pipelines, model evaluation, ETL, and Snowflake reporting.

TECHNICAL SKILLS

Programming & Data: Python, SQL, R, Pandas, NumPy, SciPy

Machine Learning: Regression, Classification, Clustering, Time-Series Forecasting, XGBoost, scikit-learn, PyTorch, Hugging Face Transformers

Statistics: Statistical Inference, Hypothesis Testing, Bootstrapping, Model Evaluation, Time-Series Backtesting

Data & Visualization: Snowflake, PostgreSQL, Oracle SQL, Neo4j, ETL, Data Validation, Power BI, matplotlib, ggplot2

Engineering: Java, Spring Boot, REST APIs, Jenkins, CI/CD, Git, Docker, Kubernetes

PROFESSIONAL EXPERIENCE

Data Analyst Intern

August 2025 – June 2026

Made By Gather | New York, NY (Remote from California)

- Built retailer-level weekly sales forecasting models in Python using Holt-Winters, ARIMA, Random Forest, and XGBoost, improving forecast accuracy from 47% to 71% and reducing MAPE by 30% on 60–150 weeks of history per retailer.
- Extended the modeling stack to item-level forecasts covering 253 Amazon SKUs, including multi-week predictions used by Demand Planning.
- Automated weekly training, model selection, and forecast generation with Python, Groovy, and Jenkins, cutting manual reporting effort from 4 hours to 1 hour.
- Reverse-engineered NetSuite sales and finance workflows into Snowflake SQL; identified join-level duplication that inflated sales metrics by more than 2x and redesigned aggregate-then-join logic for trusted Finance reporting.
- Built Power BI sales dashboards to replace manual Excel reporting, enabling stakeholders to monitor actuals, prior-year performance, and sales outlook.

Associate Software Developer

July 2022 – August 2024

Oracle Financial Services Software Ltd. | Bangalore, India

- Developed 30+ backend services and REST APIs on Oracle Banking Architecture (OBA) using Java, Spring Boot, and Oracle SQL; delivered Stop Payment Enquiry (20% fewer payment errors).
- Wrote JUnit tests for 8 services (coverage above 80%), automated Jenkins/Tomcat deployments across 40+ services, resolved production issues with SQL and logs, and mentored 4 employees and 3 interns.

PROJECTS

Motherese: Prosody-Aware Language Modeling with Whisper

April 2025 – June 2025

PyTorch, Hugging Face Transformers, OpenAI Whisper, Kubernetes, Montreal Forced Aligner, BLiMP

- Implemented a 1.55B-parameter Whisper variant in PyTorch, aligned 13 prosodic features across ~205 hours of audio, and trained YAML Kubernetes GPU jobs on UCSD Nautilus.
- Improved dependency-parsing F1 by 35% overall, with up to 66% F1 gains on hard long-range CHILDES/BLiMP tasks.

Smart Doctor and Hospital Finder

February 2025 – March 2025

Python, PostgreSQL, Neo4j, Cypher, ETL

- Built a healthcare search application over 3.2M+ Medi-Cal records with Python ETL, PostgreSQL, Neo4j, and Dijkstra specialty routing.

EDUCATION

University of California, San Diego, CA

September 2024 – June 2026

Master of Science in Data Science | GPA: 3.89/4.0

Veermata Jijabai Technological Institute (VJTI), Mumbai, India

August 2018 – May 2022

B.Tech. in Electronics and Telecommunications Engineering | CPI: 8.72/10.0